
✓ Activity 1.2 : Training Neural Networks

Objective(s):

This activity aims to demonstrate how to train neural networks using keras

Intended Learning Outcomes (ILOs):

- Demonstrate how to build and train neural networks
- Demonstrate how to evaluate and plot the model using training and validation loss

Resources:

- Jupyter Notebook

CI Pima Diabetes Dataset

- pima-indians-diabetes.csv

✓ Procedures

Load the necessary libraries

```
from __future__ import absolute_import, division, print_function

import warnings
warnings.filterwarnings("ignore")

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import confusion_matrix, precision_recall_curve, roc_auc_score, roc_curve
from sklearn.ensemble import RandomForestClassifier

import seaborn as sns

%matplotlib inline
```

```
## Import Keras objects for Deep Learning
```

```
from keras.models import Sequential
from keras.layers import Input, Dense, Flatten, Dropout, BatchNormalization
from keras.optimizers import Adam, SGD, RMSprop
```

Load the dataset

```
filepath = "pima-indians-diabetes.csv"
names = ["times_pregnant", "glucose_tolerance_test", "blood_pressure", "skin_thickness", "ir
        "bmi", "pedigree_function", "age", "has_diabetes"]
diabetes_df = pd.read_csv(filepath, names=names)
```

Check the top 5 samples of the data

```
print(diabetes_df.shape)
diabetes_df.sample(5)
```

(768, 9)

	times_pregnant	glucose_tolerance_test	blood_pressure	skin_thickness	insulin	b
662	8	167	106	46	231	37
324	2	112	75	32	0	35
494	3	80	0	0	0	0
278	5	114	74	0	0	24
353	1	90	62	12	43	27

```
diabetes_df.dtypes
```

```
times_pregnant      int64
glucose_tolerance_test  int64
blood_pressure      int64
skin_thickness      int64
insulin             int64
bmi                 float64
pedigree_function   float64
age                 int64
has_diabetes        int64
dtype: object
```

```
X = diabetes_df.iloc[:, :-1].values
y = diabetes_df["has_diabetes"].values
```

Split the data to Train, and Test (75%, 25%)

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_state=11111)

np.mean(y), np.mean(1-y)

(0.3489583333333333, 0.6510416666666666)
```

Build a single hidden layer neural network using 12 nodes. Use the sequential model with single layer network and input shape to 8.

Normalize the data

```
normalizer = StandardScaler()
X_train_norm = normalizer.fit_transform(X_train)
X_test_norm = normalizer.transform(X_test)
```

Define the model:

- Input size is 8-dimensional
- 1 hidden layer, 12 hidden nodes, sigmoid activation
- Final layer with one node and sigmoid activation (standard for binary classification)

```
model_1 = Sequential([
    Dense(12, input_shape=(8,)), activation="relu"),
    Dense(1, activation="sigmoid")
])
```

View the model summary

```
model_1.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 12)	108
dense_1 (Dense)	(None, 1)	13

Total params: 121 (484.00 Byte)
Trainable params: 121 (484.00 Byte)
Non-trainable params: 0 (0.00 Byte)

Train the model

- Compile the model with optimizer, loss function and metrics
- Use the fit function to return the run history.

```
model_1.compile(SGD(lr = .003), "binary_crossentropy", metrics=["accuracy"])  
run_hist_1 = model_1.fit(X_train_norm, y_train, validation_data=(X_test_norm, y_test), epoch
```

```
WARNING:absl:`lr` is deprecated in Keras optimizer, please use `learning_rate` or use  
Epoch 1/200  
18/18 [=====] - 1s 16ms/step - loss: 0.7618 - accuracy: 0.529  
Epoch 2/200  
18/18 [=====] - 0s 4ms/step - loss: 0.7237 - accuracy: 0.5712  
Epoch 3/200  
18/18 [=====] - 0s 4ms/step - loss: 0.6941 - accuracy: 0.6024  
Epoch 4/200  
18/18 [=====] - 0s 5ms/step - loss: 0.6702 - accuracy: 0.6181  
Epoch 5/200  
18/18 [=====] - 0s 6ms/step - loss: 0.6507 - accuracy: 0.6181  
Epoch 6/200  
18/18 [=====] - 0s 6ms/step - loss: 0.6343 - accuracy: 0.6215  
Epoch 7/200  
18/18 [=====] - 0s 5ms/step - loss: 0.6203 - accuracy: 0.6285  
Epoch 8/200  
18/18 [=====] - 0s 5ms/step - loss: 0.6083 - accuracy: 0.6458  
Epoch 9/200  
18/18 [=====] - 0s 5ms/step - loss: 0.5976 - accuracy: 0.6545  
Epoch 10/200  
18/18 [=====] - 0s 6ms/step - loss: 0.5884 - accuracy: 0.6667  
Epoch 11/200  
18/18 [=====] - 0s 5ms/step - loss: 0.5800 - accuracy: 0.6719  
Epoch 12/200  
18/18 [=====] - 0s 6ms/step - loss: 0.5725 - accuracy: 0.6788  
Epoch 13/200  
18/18 [=====] - 0s 5ms/step - loss: 0.5657 - accuracy: 0.6840  
Epoch 14/200  
18/18 [=====] - 0s 5ms/step - loss: 0.5594 - accuracy: 0.6858  
Epoch 15/200  
18/18 [=====] - 0s 6ms/step - loss: 0.5537 - accuracy: 0.6927  
Epoch 16/200  
18/18 [=====] - 0s 5ms/step - loss: 0.5485 - accuracy: 0.6944  
Epoch 17/200  
18/18 [=====] - 0s 5ms/step - loss: 0.5436 - accuracy: 0.6962  
Epoch 18/200  
18/18 [=====] - 0s 4ms/step - loss: 0.5391 - accuracy: 0.6997  
Epoch 19/200  
18/18 [=====] - 0s 5ms/step - loss: 0.5348 - accuracy: 0.7049
```

```

Epoch 20/200
18/18 [=====] - 0s 8ms/step - loss: 0.5309 - accuracy: 0.7049
Epoch 21/200
18/18 [=====] - 0s 6ms/step - loss: 0.5271 - accuracy: 0.7083
Epoch 22/200
18/18 [=====] - 0s 6ms/step - loss: 0.5238 - accuracy: 0.7066
Epoch 23/200
18/18 [=====] - 0s 6ms/step - loss: 0.5206 - accuracy: 0.7066
Epoch 24/200
18/18 [=====] - 0s 6ms/step - loss: 0.5173 - accuracy: 0.7066
Epoch 25/200
18/18 [=====] - 0s 5ms/step - loss: 0.5145 - accuracy: 0.7083
Epoch 26/200
18/18 [=====] - 0s 5ms/step - loss: 0.5116 - accuracy: 0.7118
Epoch 27/200
18/18 [=====] - 0s 6ms/step - loss: 0.5090 - accuracy: 0.7205
Epoch 28/200
18/18 [=====] - 0s 5ms/step - loss: 0.5064 - accuracy: 0.7257

```

```

## Like we did for the Random Forest, we generate two kinds of predictions
# One is a hard decision, the other is a probabilistic score.

```

```

y_pred_class_nn_1 = model_1.predict(X_test_norm)
y_pred_prob_nn_1 = model_1.predict(X_test_norm)

```

```

6/6 [=====] - 0s 2ms/step
6/6 [=====] - 0s 2ms/step

```

```

# Let's check out the outputs to get a feel for how keras apis work.
y_pred_class_nn_1[:10]

```

```

array([[0.65582037],
       [0.6507114 ],
       [0.32304847],
       [0.21720952],
       [0.1633364 ],
       [0.4551786 ],
       [0.02752979],
       [0.38203385],
       [0.89780146],
       [0.17257361]], dtype=float32)

```

```

y_pred_prob_nn_1[:10]

```

```

array([[0.65582037],
       [0.6507114 ],
       [0.32304847],
       [0.21720952],
       [0.1633364 ],
       [0.4551786 ],
       [0.02752979],
       [0.38203385],

```

```
[0.89780146],  
[0.17257361]], dtype=float32)
```

Create the `plot_roc` function

```
rf = RandomForestClassifier(n_estimators=200)  
rf.fit(X_train, y_train)
```

```
▼      RandomForestClassifier  
RandomForestClassifier(n_estimators=200)
```

```
def plot_roc(y_test, y_pred, model_name):  
    fpr, tpr, thr = roc_curve(y_test, y_pred)  
    fig, ax = plt.subplots(figsize=(8, 8))  
    ax.plot(fpr, tpr, 'k-')  
    ax.plot([0, 1], [0, 1], 'k--', linewidth=.5) # roc curve for random model  
    ax.grid(True)  
    ax.set(title='ROC Curve for {} on PIMA diabetes problem'.format(model_name),  
           xlim=[-0.01, 1.01], ylim=[-0.01, 1.01])
```

Evaluate the model performance and plot the ROC CURVE

```
y_pred_class_rf = rf.predict(X_test)  
y_pred_prob_rf = rf.predict_proba(X_test)
```


```
print('accuracy is {:.3f}'.format(accuracy_score(y_test, y_pred_class_rf)))  
print('roc-auc is {:.3f}'.format(roc_auc_score(y_test, y_pred_prob_rf[:,1])))
```

```
accuracy is 0.760  
roc-auc is 0.818
```

```
print('accuracy is {:.3f}'.format(accuracy_score(y_test, y_pred_class_rf)))  
print('roc-auc is {:.3f}'.format(roc_auc_score(y_test, y_pred_prob_rf[:,1])))
```

```
plot_roc(y_test, y_pred_prob_nn_1, 'NN')
```

```
accuracy is 0.760  
roc-auc is 0.818
```




Plot the training loss and the validation loss over the different epochs and see how it looks

```
run_hist_1.history.keys()  
  
dict_keys(['loss', 'accuracy', 'val_loss', 'val_accuracy'])
```

```
fig, ax = plt.subplots()
ax.plot(run_hist_1.history["loss"], 'r', marker='.', label="Train Loss")
ax.plot(run_hist_1.history["val_loss"], 'b', marker='.', label="Validation Loss")
ax.legend()
```

<matplotlib.legend.Legend at 0x7ce35e04a7a0>



What is your interpretation about the result of the train and validation loss?

✓ type your answer here

My interpretation on the results of train and validation loss is that the model should have further training, thus adding more epoch may result into better or worse. Also, the results shows it is underfitting since the both losses is high creating wider gap each epoch.

✓ Supplementary Activity

- Build a model with two hidden layers, each with 6 nodes
- Use the "relu" activation function for the hidden layers, and "sigmoid" for the final layer
- Use a learning rate of .003 and train for 1500 epochs
- Graph the trajectory of the loss functions, accuracy on both train and test set

- Plot the roc curve for the predictions
- Use different learning rates, numbers of epochs, and network structures.
- Plot the results of training and validation loss using different learning rates, number of epochs and network structures
- Interpret your result

```
#Build a model with two hidden layers, each with 6 nodes
#Use the "relu" activation function for the hidden layers, and "sigmoid" for the final layer
model_supple = Sequential([
    Dense(6, input_shape=(8,)), activation="sigmoid"),
    Dense(2, activation="relu")
])
```

```
model_supple.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
dense_2 (Dense)	(None, 6)	54
dense_3 (Dense)	(None, 2)	14
Total params: 68 (272.00 Byte)		
Trainable params: 68 (272.00 Byte)		
Non-trainable params: 0 (0.00 Byte)		

```
model_supple.compile(SGD(lr = .003), "binary_crossentropy", metrics=["accuracy"])
run_hist_1 = model_supple.fit(X_train_norm, y_train, validation_data=(X_test_norm, y_test),
```

```
WARNING:absl:`lr` is deprecated in Keras optimizer, please use `learning_rate` or use
Epoch 1/1500
18/18 [=====] - 1s 13ms/step - loss: 6.6705 - accuracy: 0.6545
Epoch 2/1500
18/18 [=====] - 0s 4ms/step - loss: 6.0002 - accuracy: 0.6545
Epoch 3/1500
18/18 [=====] - 0s 4ms/step - loss: 5.5249 - accuracy: 0.6528
Epoch 4/1500
18/18 [=====] - 0s 4ms/step - loss: 3.2684 - accuracy: 0.6545
Epoch 5/1500
18/18 [=====] - 0s 4ms/step - loss: 3.2098 - accuracy: 0.6545
Epoch 6/1500
18/18 [=====] - 0s 5ms/step - loss: 3.0724 - accuracy: 0.6545
Epoch 7/1500
18/18 [=====] - 0s 5ms/step - loss: 3.0379 - accuracy: 0.6545
Epoch 8/1500
18/18 [=====] - 0s 4ms/step - loss: 3.0258 - accuracy: 0.6545
Epoch 9/1500
18/18 [=====] - 0s 4ms/step - loss: 3.0175 - accuracy: 0.6545
```

```

Epoch 10/1500
18/18 [=====] - 0s 4ms/step - loss: 2.9997 - accuracy: 0.6545
Epoch 11/1500
18/18 [=====] - 0s 5ms/step - loss: 2.9832 - accuracy: 0.6545
Epoch 12/1500
18/18 [=====] - 0s 4ms/step - loss: 2.9775 - accuracy: 0.6545
Epoch 13/1500
18/18 [=====] - 0s 4ms/step - loss: 2.9723 - accuracy: 0.6545
Epoch 14/1500
18/18 [=====] - 0s 4ms/step - loss: 2.9684 - accuracy: 0.6545
Epoch 15/1500
18/18 [=====] - 0s 4ms/step - loss: 2.9649 - accuracy: 0.6545
Epoch 16/1500
18/18 [=====] - 0s 3ms/step - loss: 2.9615 - accuracy: 0.6545
Epoch 17/1500
18/18 [=====] - 0s 4ms/step - loss: 2.9589 - accuracy: 0.6545
Epoch 18/1500
18/18 [=====] - 0s 4ms/step - loss: 2.9567 - accuracy: 0.6545
Epoch 19/1500
18/18 [=====] - 0s 4ms/step - loss: 2.9541 - accuracy: 0.6545
Epoch 20/1500
18/18 [=====] - 0s 4ms/step - loss: 2.9524 - accuracy: 0.6545
Epoch 21/1500
18/18 [=====] - 0s 4ms/step - loss: 2.9504 - accuracy: 0.6545
Epoch 22/1500
18/18 [=====] - 0s 4ms/step - loss: 2.9487 - accuracy: 0.6545
Epoch 23/1500
18/18 [=====] - 0s 5ms/step - loss: 2.9472 - accuracy: 0.6545
Epoch 24/1500
18/18 [=====] - 0s 4ms/step - loss: 2.9456 - accuracy: 0.6545
Epoch 25/1500
18/18 [=====] - 0s 4ms/step - loss: 3.0134 - accuracy: 0.6545
Epoch 26/1500
18/18 [=====] - 0s 4ms/step - loss: 2.9322 - accuracy: 0.6545
Epoch 27/1500
18/18 [=====] - 0s 4ms/step - loss: 2.9265 - accuracy: 0.6545
Epoch 28/1500
18/18 [=====] - 0s 4ms/step - loss: 2.9249 - accuracy: 0.6545
Epoch 29/1500
18/18 [=====] - 0s 4ms/step - loss: 2.9249 - accuracy: 0.6545

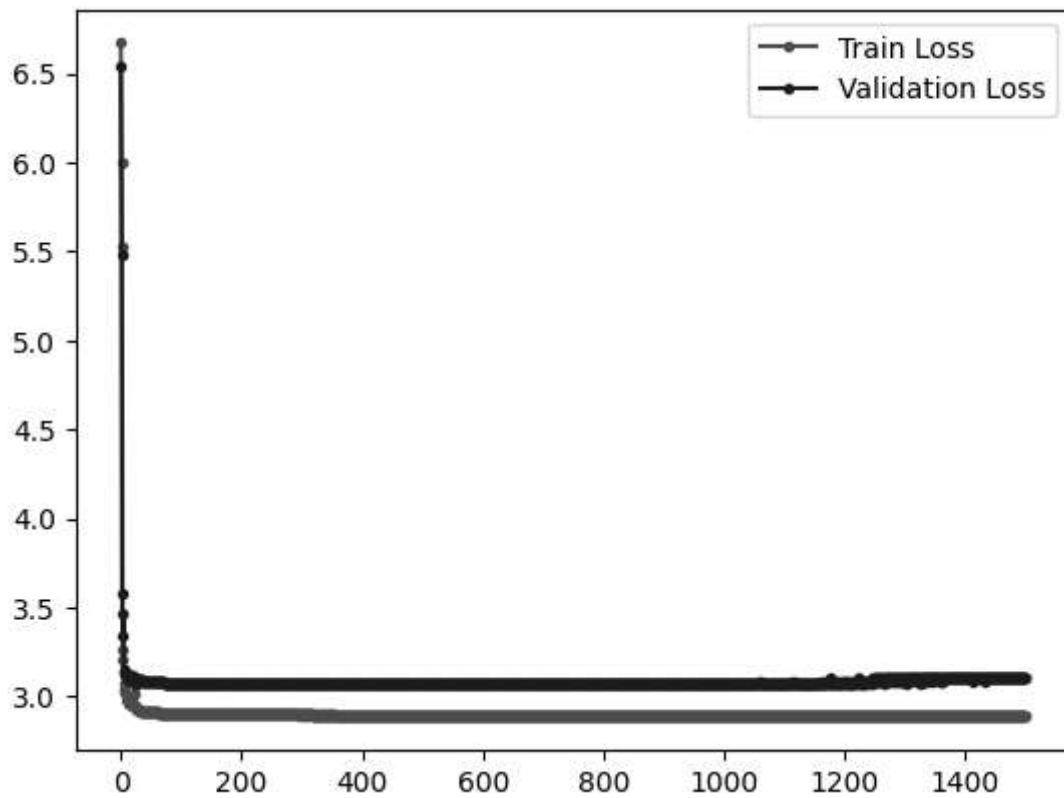
```

```

fig, ax = plt.subplots()
ax.plot(run_hist_1.history["loss"], 'r', marker='.', label="Train Loss")
ax.plot(run_hist_1.history["val_loss"], 'b', marker='.', label="Validation Loss")
ax.legend()

```

<matplotlib.legend.Legend at 0x7ce35defc6a0>




- Interpret results The results shows in the graph that the model is underfitting since the gap between the two losses are high. It shows that under 200 epoch the model are good fit indicating that having 200 epoch is enough.

```
#Plot the roc curve for the predictions
print('accuracy is {:.3f}'.format(accuracy_score(y_test,y_pred_class_rf)))
print('roc-auc is {:.3f}'.format(roc_auc_score(y_test,y_pred_prob_rf[:,1])))

plot_roc(y_test, y_pred_prob_nn_1, 'NN')
```

accuracy is 0.760
roc-auc is 0.818



```
#Use different learning rates, numbers of epochs, and network structures.
model_2 = Sequential([
    Dense(10, input_shape=(8,), activation="tanh"),
    Dense(1, activation="sigmoid")
])
```

```
model_2.summary()
```

Model: "sequential_2"

Layer (type)	Output Shape	Param #
--------------	--------------	---------

```
=====
dense_4 (Dense)              (None, 10)              90

dense_5 (Dense)              (None, 1)               11

=====
Total params: 101 (404.00 Byte)
Trainable params: 101 (404.00 Byte)
Non-trainable params: 0 (0.00 Byte)
=====
```

```
#Plot the results of training and validation loss using different learning rates, number of
model_2.compile(SGD(lr = .005), "binary_crossentropy", metrics=["accuracy"])
run_hist_2 = model_2.fit(X_train_norm, y_train, validation_data=(X_test_norm, y_test), epoch
```

```
WARNING:absl:`lr` is deprecated in Keras optimizer, please use `learning_rate` or use ▲
Epoch 1/500
18/18 [=====] - 1s 13ms/step - loss: 0.6533 - accuracy: 0.631
Epoch 2/500
18/18 [=====] - 0s 5ms/step - loss: 0.6307 - accuracy: 0.6806
Epoch 3/500
18/18 [=====] - 0s 4ms/step - loss: 0.6120 - accuracy: 0.7101
Epoch 4/500
18/18 [=====] - 0s 5ms/step - loss: 0.5962 - accuracy: 0.7222
Epoch 5/500
18/18 [=====] - 0s 4ms/step - loss: 0.5831 - accuracy: 0.7361
Epoch 6/500
18/18 [=====] - 0s 4ms/step - loss: 0.5716 - accuracy: 0.7344
Epoch 7/500
18/18 [=====] - 0s 4ms/step - loss: 0.5618 - accuracy: 0.7413
Epoch 8/500
18/18 [=====] - 0s 4ms/step - loss: 0.5533 - accuracy: 0.7448
Epoch 9/500
18/18 [=====] - 0s 4ms/step - loss: 0.5458 - accuracy: 0.7506
Epoch 10/500
18/18 [=====] - 0s 4ms/step - loss: 0.5392 - accuracy: 0.7517
Epoch 11/500
18/18 [=====] - 0s 4ms/step - loss: 0.5334 - accuracy: 0.7483
Epoch 12/500
18/18 [=====] - 0s 4ms/step - loss: 0.5282 - accuracy: 0.7483
Epoch 13/500
18/18 [=====] - 0s 5ms/step - loss: 0.5235 - accuracy: 0.7517
Epoch 14/500
18/18 [=====] - 0s 5ms/step - loss: 0.5194 - accuracy: 0.7604
Epoch 15/500
18/18 [=====] - 0s 4ms/step - loss: 0.5156 - accuracy: 0.7639
Epoch 16/500
18/18 [=====] - 0s 4ms/step - loss: 0.5122 - accuracy: 0.7656
Epoch 17/500
18/18 [=====] - 0s 4ms/step - loss: 0.5090 - accuracy: 0.7639
Epoch 18/500
18/18 [=====] - 0s 4ms/step - loss: 0.5061 - accuracy: 0.7656
Epoch 19/500
18/18 [=====] - 0s 4ms/step - loss: 0.5036 - accuracy: 0.7691
Epoch 20/500
```

```

18/18 [=====] - 0s 5ms/step - loss: 0.5011 - accuracy: 0.7708
Epoch 21/500
18/18 [=====] - 0s 4ms/step - loss: 0.4989 - accuracy: 0.7708
Epoch 22/500
18/18 [=====] - 0s 4ms/step - loss: 0.4968 - accuracy: 0.7708
Epoch 23/500
18/18 [=====] - 0s 4ms/step - loss: 0.4949 - accuracy: 0.7726
Epoch 24/500
18/18 [=====] - 0s 4ms/step - loss: 0.4931 - accuracy: 0.7760
Epoch 25/500
18/18 [=====] - 0s 4ms/step - loss: 0.4914 - accuracy: 0.7760
Epoch 26/500
18/18 [=====] - 0s 4ms/step - loss: 0.4899 - accuracy: 0.7778
Epoch 27/500
18/18 [=====] - 0s 4ms/step - loss: 0.4884 - accuracy: 0.7778
Epoch 28/500
18/18 [=====] - 0s 4ms/step - loss: 0.4871 - accuracy: 0.7795


```

```

fig, ax = plt.subplots()
ax.plot(run_hist_2.history["loss"], 'r', marker='.', label="Train Loss")
ax.plot(run_hist_2.history["val_loss"], 'b', marker='.', label="Validation Loss")
ax.legend()

```

<matplotlib.legend.Legend at 0x7ce35dbe9f00>



- Interpret your result

It shows that the experimented one is good fit since both are stabilizing as the epoch gets higher. The results shows stablization as the epoch gets further.

✓ Conclusion

To conclude, this activity help me understand on how to train neural networks and how to apply it into deep learning and how to use it in models. Although I experienced some errors, I still tried to solve it and I think I solved it somehow and as the time goes by and everytime I tried to fix the errors I realized what this hands-on activity for and how this works. Lastly, this activity help me to have a better idea how cool deep learning is and how important it is.

Double-click (or enter) to edit

Double-click (or enter) to edit

+ Code

+ Text